


# An Intro to Deep Learning

Emmanuel Bengio

June 6, 2017

# What is deep learning?



Credit: Graphcore/Matt Fyles

# What is deep learning?


The technical answer:

- Stacking layers of linear transforms and non-linearities


The less-technical answer:

- Stacking layers of abstractions
- Brain-like structure
- Learning **representations**


# Linear regression



# Adding intermediate nodes




# Adding intermediate nodes



(By the way with this you can theoretically represent any function!)

# Adding more layers



# In math

In terms of linear algebra, no hidden layers (linear regression):

$$\hat{y} = \mathbf{w}^T \mathbf{x}$$



# In math

In terms of linear algebra, no hidden layers (linear regression):

$$\hat{y} = \mathbf{w}^\top \mathbf{x}$$


A single hidden layer:

$$\hat{y} = \mathbf{w}^\top f(W\mathbf{x} + \mathbf{b})$$

# Activation Functions

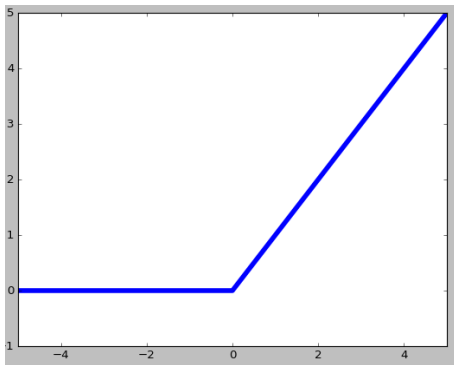
Activation functions are applied *element-wise* to each element of a vector. For example the sigmoid:

$$f(x) = \frac{1}{1 + e^{-x}}$$



Right now the most popular activation function is the Rectified Linear Unit (ReLU):

$$f(x) = \max\{0, x\}$$



# Training Neural Networks via Gradient Descent

Training NNets is conceptually rather simple:

- Show an example  $x$
- Compute  $\hat{y}$  and the error (e.g.  $E = (\hat{y} - y)^2$ )
- Adjust the weights  $w$  by  $\alpha \frac{\partial E}{\partial w}$
- Repeat

$\alpha$  is called a learning rate.

# Training Neural Networks via Gradient Descent

$$\frac{\partial E}{\partial w} \text{ or } \nabla_w E$$

is called the gradient of the error with respect to the parameters.  
Deep learning libraries compute it for you!


# Deep Learning Software

- Theano, Tensorflow, MXNet, PyTorch
- Keras, Lasagne

Theano is older but stable. Tensorflow is more recent, engineered-feeling. MXNet/PyTorch are less verbose and also more recent.

Keras and Lasagne are built on top of these libraries, abstract NNet operations.

# A sequence of blocks



# A sequence of blocks

```
import keras
from keras.models import Sequential
from keras.layers import Dense
```



## A sequence of blocks

```
model = Sequential()  
model.add(Dense(32, activation='relu',  
                input_shape=10))  
model.add(Dense(64, activation='relu'))  
model.add(Dense(10, activation='softmax'))
```

# Compiling the model

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```

# Training the model via gradient descent

```
model.fit(X_train, Y_train,  
          batch_size=32, epochs=10, verbose=1)
```

## Evaluating the model on new data

```
score = model.evaluate(X_test, Y_test, verbose=0)
```